

```
In [174... import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
from tensorflow.keras.losses import BinaryCrossentropy
import matplotlib.pyplot as plt
import os
from sklearn.metrics import classification_report, f1_score, accuracy_score, precision_score, recall_score
import numpy as np
```

```
In [175... fresh = len(os.listdir("/Users/alyaaljarallah/Desktop/Fruit_project/fruits_dataset/fresh"))
rotten = len(os.listdir("/Users/alyaaljarallah/Desktop/Fruit_project/fruits_dataset/rotten"))

print(f"Fresh: {fresh}, Rotten: {rotten}")
```

Fresh: 105, Rotten: 121

```
In [176... dataset_path = "/Users/alyaaljarallah/Desktop/Fruit_project/fruits_dataset"

train_datagen = ImageDataGenerator(
    rescale=1./255,
    validation_split=0.2,
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2, zoom_range=0.3,
    horizontal_flip=True,
    brightness_range=[0.6, 1.4]
)

val_datagen = ImageDataGenerator(
    rescale=1./255,
    validation_split=0.2
)

train_generator = train_datagen.flow_from_directory( dataset_path,
target_size=(150, 150), batch_size=32, class_mode='binary', subset='training', shuffle=True)

val_generator = val_datagen.flow_from_directory( dataset_path,
target_size=(150, 150), batch_size=32, class_mode='binary', subset='validation', shuffle=False )
```

Found 181 images belonging to 2 classes.

Found 44 images belonging to 2 classes.

In [177...

```
model = Sequential([
    Conv2D(32, (3, 3), activation='relu',
          input_shape=(150, 150, 3)), MaxPooling2D(2, 2),

    Conv2D(64, (3, 3), activation='relu'), MaxPooling2D(2, 2),

    Conv2D(128, (3, 3), activation='relu'), MaxPooling2D(2, 2),

    Flatten(),
    Dense(128, activation='relu'), Dropout(0.3),
    Dense(1, activation='sigmoid')
])

model.compile(optimizer=Adam(learning_rate=1e-4),
              loss=BinaryCrossentropy(label_smoothing=0.1), metrics=['accuracy'])
model.summary()
```

/Users/alyaaljarallah/opt/anaconda3/lib/python3.9/site-packages/keras/src/layers/convolutional/base_conv.py:107: User Warning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential_23"

Layer (type)	Output Shape	Param #
conv2d_56 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_46 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_57 (Conv2D)	(None, 72, 72, 64)	18,496
max_pooling2d_47 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_58 (Conv2D)	(None, 34, 34, 128)	73,856
max_pooling2d_48 (MaxPooling2D)	(None, 17, 17, 128)	0
flatten_23 (Flatten)	(None, 36992)	0
dense_48 (Dense)	(None, 128)	4,735,104
dropout_19 (Dropout)	(None, 128)	0
dense_49 (Dense)	(None, 1)	129

Total params: 4,828,481 (18.42 MB)

Trainable params: 4,828,481 (18.42 MB)


Non-trainable params: 0 (0.00 B)


```
In [178... early_stop = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)


checkpoint = ModelCheckpoint( 'best_model.keras', monitor='val_accuracy', save_best_only=True, mode='max')


history = model.fit( train_generator, validation_data=val_generator, epochs=20, callbacks=[early_stop, checkpoint] )
```


```
/Users/alyaaljarallah/opt/anaconda3/lib/python3.9/site-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.
self._warn_if_super_not_called()
```


Epoch 1/20
6/6  **10s** 982ms/step - accuracy: 0.5701 - loss: 0.6855 - val_accuracy: 0.5455 - val_loss: 0.6593


Epoch 2/20
6/6  **6s** 860ms/step - accuracy: 0.6048 - loss: 0.6414 - val_accuracy: 0.5455 - val_loss: 0.6544


Epoch 3/20
6/6  **12s** 1s/step - accuracy: 0.6467 - loss: 0.6385 - val_accuracy: 0.8182 - val_loss: 0.5700


Epoch 4/20
6/6  **7s** 944ms/step - accuracy: 0.6496 - loss: 0.6403 - val_accuracy: 0.6591 - val_loss: 0.6176


Epoch 5/20
6/6  **6s** 1s/step - accuracy: 0.7192 - loss: 0.6049 - val_accuracy: 0.8636 - val_loss: 0.5215


Epoch 6/20
6/6  **6s** 958ms/step - accuracy: 0.6688 - loss: 0.6110 - val_accuracy: 0.7273 - val_loss: 0.5704


Epoch 7/20
6/6  **7s** 1s/step - accuracy: 0.6960 - loss: 0.5845 - val_accuracy: 0.9318 - val_loss: 0.5035


Epoch 8/20
6/6  **7s** 1s/step - accuracy: 0.7139 - loss: 0.5537 - val_accuracy: 0.9091 - val_loss: 0.4950


Epoch 9/20
6/6  **9s** 881ms/step - accuracy: 0.7741 - loss: 0.5568 - val_accuracy: 0.9091 - val_loss: 0.4440


Epoch 10/20
6/6  **6s** 1s/step - accuracy: 0.7679 - loss: 0.5398 - val_accuracy: 0.9545 - val_loss: 0.4471


Epoch 11/20
6/6  **7s** 1s/step - accuracy: 0.7532 - loss: 0.5659 - val_accuracy: 0.9773 - val_loss: 0.4105


Epoch 12/20
6/6  **5s** 808ms/step - accuracy: 0.7827 - loss: 0.5110 - val_accuracy: 0.9545 - val_loss: 0.3975


Epoch 13/20
6/6  **6s** 912ms/step - accuracy: 0.8088 - loss: 0.5162 - val_accuracy: 0.8864 - val_loss: 0.4228


Epoch 14/20
6/6  **6s** 970ms/step - accuracy: 0.8012 - loss: 0.5212 - val_accuracy: 0.9545 - val_loss: 0.3943


Epoch 15/20
6/6  **7s** 962ms/step - accuracy: 0.8032 - loss: 0.5300 - val_accuracy: 0.9091 - val_loss: 0.3823

Epoch 16/20
6/6  **7s** 1s/step - accuracy: 0.8314 - loss: 0.4487 - val_accuracy: 0.9545 - val_loss: 0.3741

Epoch 17/20
6/6  **9s** 854ms/step - accuracy: 0.8347 - loss: 0.4897 - val_accuracy: 0.9318 - val_loss: 0.3561

Epoch 18/20
6/6  **5s** 707ms/step - accuracy: 0.8288 - loss: 0.4550 - val_accuracy: 0.9545 - val_loss: 0.3242

Epoch 19/20
6/6  **4s** 689ms/step - accuracy: 0.8223 - loss: 0.4702 - val_accuracy: 0.9545 - val_loss: 0.3164

Epoch 20/20
6/6  **4s** 738ms/step - accuracy: 0.8680 - loss: 0.4463 - val_accuracy: 0.9545 - val_loss: 0.3315

```
In [181... val_generator.reset()
y_probs = model.predict(val_generator)
y_preds = (y_probs > 0.5).astype(int).flatten()
```

```
y_true = val_generator.classes[:len(y_preds)]

accuracy = accuracy_score(y_true, y_preds)
precision = precision_score(y_true, y_preds)
recall = recall_score(y_true, y_preds)
f1 = f1_score(y_true, y_preds)

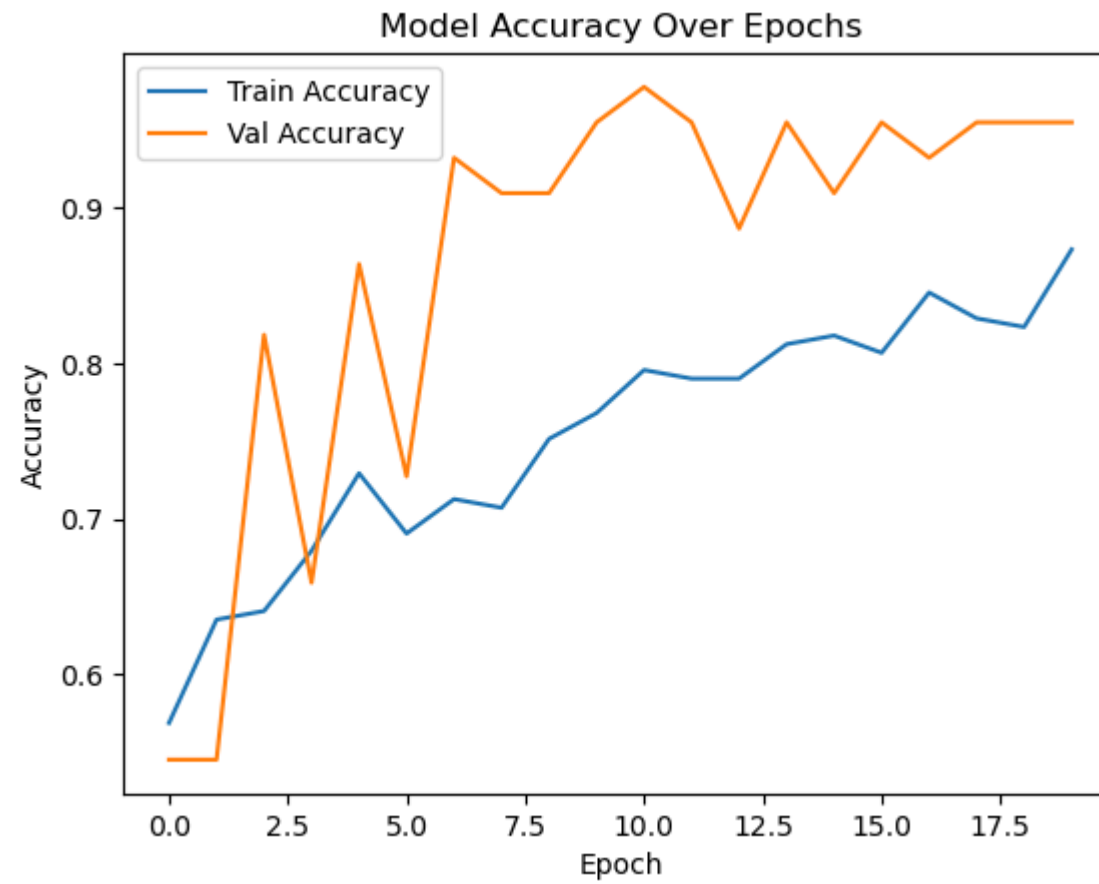
print(f"Accuracy:  {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall:    {recall:.4f}")
print(f"F1 Score:   {f1:.4f}")

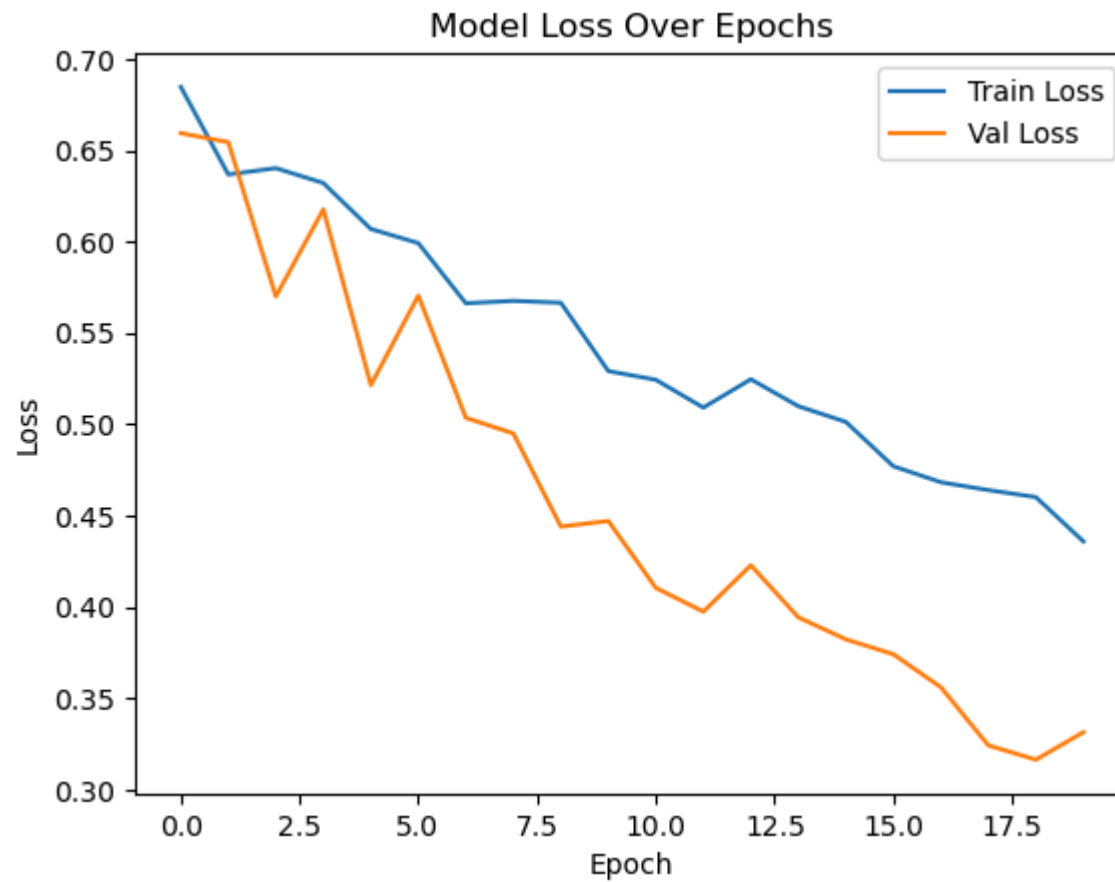
# Accuracy
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Val Accuracy')
plt.title('Model Accuracy Over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

# Loss
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title('Model Loss Over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

2/2 ————— 1s 152ms/step

Accuracy: 0.9545
Precision: 1.0000
Recall: 0.9167
F1 Score: 0.9565





```
In [184... plt.figure(figsize=(10, 6))
for i in range(6):
    plt.subplot(2, 3, i + 1)
    plt.imshow(x_val[i])
    plt.title("Fresh" if y_val[i] == 0 else "Rotten")
    plt.axis('off')

plt.tight_layout()
plt.show()
```

Fresh



Fresh



Fresh



Fresh



Fresh



Fresh

